

# Programming Assignment 2

Simon Bolding  
CSCE 626

March 29, 2015

# References

1. People in class - Hans Hammer, Daniel Holladay, Josh Hansel
2. [wikipedia.org/wiki/Prefix\\_sum](http://wikipedia.org/wiki/Prefix_sum), [computing.llnl.gov](http://computing.llnl.gov), [mpitutorial.com](http://mpitutorial.com)
3. Class notes
4. EOS website: [sc.tamu.edu/systems/eos](http://sc.tamu.edu/systems/eos)

## 1 Introduction

In this work experimental analysis was performed to evaluate the performance of two different methods for computing prefix sums in parallel: algorithm implemented with MPI and an algorithm implemented using OpenMP. Given a sequence of numbers  $x_1, x_2, \dots, x_n$ , the prefix sums are the partial sums

$$\begin{aligned}s_1 &= x_1 \\ s_2 &= x_1 + x_2 \\ &\vdots \\ s_n &= x_1 + x_2 + \dots + x_n.\end{aligned}$$

Multiple experiments were performed, and the results were analyzed to compare the performance of the algorithms as a function of the number of parallel processors used, for both algorithms. Coefficients for the theoretical asymptotic complexity were determined. Both strong and weak scaling studies were performed as well, and speed up studied as a function of the number of elements. The algorithms are benchmarked against the best sequential algorithm.

## 2 Theoretical Analysis

### 2.1 Sequential Algorithm

An efficient, time optimal sequential algorithm was used as a reference to compare to parallel algorithms. The algorithm is given in Alg. 1. This algorithm has asymptotic time complexity  $T(n) = O(n)$ , where  $n$  is the number of input integers. Because  $n - 1$  additions must be performed to calculate the prefix sum, this is clearly the optimal sequential algorithm. It is noted as an implementation detail that the call to `PrefixSumSerial` uses the same function call in the parallel versions of the code. This ensures there are no optimization differences for this function call.

### 2.2 MPI Algorithm

The algorithm for the MPI implementation is given in Alg. 2. This algorithm uses the standard approach from the previous homework of computing local prefix sums, followed by

---

**Algorithm 1** Serial prefix sum algorithm

---

```
procedure PREFIXSUMSERIAL( $A, B$ ) ▷ Assumes  $\text{len}(A)=\text{len}(B)=n$ 
  variables:  $i, n$ 
   $n = \text{len}(A)$ 
   $B(0) := A(0)$  ▷ Initialize prefix sum
  for  $i = 1$  to  $n - 1$  do
     $B(i) := A(i) + B(i-1)$ 
  end for ▷  $B$  now contains the prefix sum of  $A$ 
end procedure
```

---

a global prefix sum on the last element in each local array, before adding the sum of previous processors to each element.

The algorithm assumes that the input elements and output prefix elements are distributed across  $p$  processors. Thus, for  $p$  processors and  $n$  integers, each processor contains  $m \approx n/p$  elements. If  $n/p$  is not an integer, then the remaining elements are distributed, one per processor, amongst the first  $n \bmod p$  elements. Each processor computes the prefix sum of its  $m$  elements, and order  $O(n/p)$  operation. Using `MPI_Allgather`, a copy of the last element in each processor's prefix sum array (representing the sum of all elements on each processor) is copied to all processors. Thus, each processor now contains an array  $C$  of size  $p$ , with the  $i$ -th element containing the sum of the elements on the  $i$ -th processor. The prefix sum of  $C$  is computed by each processor, an  $O(p)$  operation. Finally, each processor adds the sum of integers from all previous processors to its local prefix sum array as necessary. each processors values.

For time and work complexity, the initial prefix sums and final value additions require  $O(n/p)$  time and  $O(n)$  work. To perform the global prefix sum, the `MPI_Allgather` command was utilized to copy the processor sums to each processor. Although  $p$  numbers must be communicated, by combining messages in a tree-reduction, this command only requires the large communication overhead of  $O(\log p)$  messages. Thus, as far as asymptotic complexity, this process represents an  $O(\log p)$  communication time, with  $p \log p$  work. The serial prefix sum of each processors array is computed only up to that processor, but this is bounded by an  $O(p)$  time and  $O(p^2)$  work. Thus, in total,  $T(n) = O(n/p + p + \log(p)) = O(n/p + p)$  and  $W(n) = O(n + p^2 + p \log p) = O(n + p^2)$ . It is noted that for the case of the experiments performed in this work,  $n \gg p$ , and the  $O(p)$  terms become negligible. We have verified this algorithm works in previous homework.

## 2.3 OpenMP Algorithm

The OpenMP algorithm is given in Alg. 3. The OpenMP algorithm works similarly to the MPI algorithm above. Each thread performs the local prefix sum on its portion of the input array, before a serial prefix sum is performed on the  $p$  processor sums on an additional array. This step is performed with the simple serial algorithm by one processor. This is done because for the shared memory experiments performed,  $p \ll n$ , and the memory overhead of using a more complicated algorithm would ultimately result in a less efficient algorithm

---

**Algorithm 2** MPI prefix sum algorithm

---

**procedure** PREFIXSUMMPI( $A$ ) ▷ Assumes  $\text{len}(A)=n$   
    ▷ Assume  $A$  and  $B$  distributed evenly on processors,  $\approx n/p$  to each processor  
    Local variables:  $id, p, m$  ▷ MPI Rank, number of processors, my size  
    Local variables:  $C, D$  ▷ Arrays of size  $p$  for parallel prefix sum  
     $id := \text{get\_my\_id}()$  ▷ Procs. numbered  $0, 1, \dots, n-1$   
     $m := \lfloor n/p \rfloor$  ▷ Determine size on each processor  
    **if**  $id < n \bmod p$  **then**  
         $m := m + 1$   
    **end if**  
    Local variables:  $myints$  ▷ this procs'  $m$  input values of  $A$   
    Local variables:  $mypsums$  ▷ this procs'  $m$  prefix sum values of  $B$   
    **if**  $\text{len}(A) = 1$  **then**  
         $B(0) := A(0)$   
        **return**  
    **else**  
        PrefixSumSerial( $myints, mypsums$ )  
    **end if**  
    MPIAllgather( $mypsums(m-1), C$ ) ▷  $C$  now contains copy of sum of elements for  
    each of  $p$  processors  
    **if**  $id > 0$  **then**  $D(0) := C(0)$   
        **for**  $i = 0$  to  $id - 1$  **do** ▷ All Procs perform prefix sum on their local copy of  $C$   
             $D(i) := D(i-1) + C(i)$   
        **end for**  
        **for**  $i = 0$  to  $n$  **do**  
             $mypsums(i) := mypsums(i) + D(id-1)$   
        **end for**  
    **end if** ▷ Collectively the  $p$  arrays  $mypsums$  contain prefix sum of  $A$   
**end procedure**

---

run time. Finally, the sum of previous processor's values are added to each portion of the array in parallel, as necessary.

For time complexity, the initial prefix sum and addition of values at the end are  $O(n/p)$ , with work complexity  $O(n)$ . The prefix sum of the processor sums is  $O(p)$  time, but, because there is shared memory access, this step can be performed by a single processor. This results in  $O(p)$  work. In total, for this algorithm,  $T(n) = O(n/p + p)$ , where  $p \ll n$ , and Work is  $O(n + p)$ . We have verified this algorithm works in previous homework.

---

**Algorithm 3** OpenMP prefix sum algorithm

---

```

procedure PREFIXSUMOPENMP( $A$ )                                ▷ Assumes  $\text{len}(A)=n$ 
    ▷ Assume  $A$  and  $B$  distributed evenly on processors,  $\approx n/p$  to each processor
    Local variables:  $id, p, m$                                 ▷ MPI Rank, number of processors, my size
    Local variables:  $C, D$                                     ▷ Arrays of size  $p$  for parallel prefix sum
     $id := \text{get\_my\_id}()$                                     ▷ Procs. numbered  $0, 1, \dots, n-1$ 
     $m := \lfloor n/p \rfloor$                                         ▷ Determine size on each processor
    if  $id < n \bmod p$  then
         $m := m + 1$ 
    end if
    Local variables:  $myints$                                 ▷ this procs'  $m$  input values of  $A$ 
    Local variables:  $mypsums$                                 ▷ this procs'  $m$  prefix sum values of  $B$ 
    if  $\text{len}(A) = 1$  then
         $B(0) := A(0)$ 
        return
    else
        PrefixSumSerial( $myints, mypsums$ )
    end if
    MPIAllgather( $mypsums(m-1), C$ )    ▷  $C$  now contains copy of sum of elements for
    each of  $p$  processors
    if  $id > 0$  then  $D(0) := C(0)$ 
        for  $i = 0$  to  $id - 1$  do    ▷ All Procs perform prefix sum on their local copy of  $C$ 
             $D(i) := D(i-1) + C(i)$ 
        end for
        for  $i = 0$  to  $n$  do
             $mypsums(i) := mypsums(i) + D(id-1)$ 
        end for
    end if    ▷ Collectively the  $p$  arrays  $mypsums$  contain prefix sum of  $A$ 
end procedure

```

---

## 3 Experimental Setup

### 3.1 Machine Information

The parallel programs were tested on *eos*, a machine at Texas A&M. For all the results in this work, the available Intel “Nehalem” nodes were used. These processors use Intel 64-bit architecture. Each node contains two sockets, each attached to 4 processors, resulting in 8 processing units per core. There is some potential difference in memory access times on the chip when going from 4 to 8 cores, where memory must be accessed off chip. The nodes are connected in a “Fat Tree” topology. This results in a constant communication time to access any off board node from any other.

For memory, each core has 32 kB L1 cache and 256 kB of unified L2 cache. Each Nehalem chip (containing four cores) has an 8 MB shared L3 cache. There is  $\sim 22$  GB of shared RAM on each node (2 chips, or 8 cores). The RAM has non-uniform access time, with longer access times when a core accesses the DRAM that is not near it. The RAM has non-uniform access time, with longer access times when a core accesses the DRAM that is not near it. More details about the architecture on *eos* can be found at <http://sc.tamu.edu/systems/eos/>

### 3.2 Description of Experiments

Several experiments were performed to gauge the efficiency of the Alg. 2 and 3. They are discussed individually below. Batch files to run the various jobs were created using a Python script. Output files were processed with a Python script as well. Example scripts for the strong scaling study were included in the submitted tar ball.

#### 3.2.1 Strong Scaling Study

The purpose of the strong scaling study is to see how much faster a problem of a fixed size can be solved by using more processors. Speed up was used as a performance measure for the strong scaling study. Speed up is defined as the ratio  $T_{ser}/T_p$ , where  $T_{ser}$  is the time to solve the problem using the most efficient serial algorithm and  $T_p$  is the time to solve the program using  $p$  processors with a particular parallel algorithm. This is different than scalability, which is the ratio  $T_1/T_p$ , which can also be used as a performance measure for strong scaling. A program which scales perfectly would show a linear, one-to-one speed up. In general this is not the case due to communication and memory overhead.

For both the MPI and OpenMP algorithm,  $10^9$  integers was chosen as the input problem size. The choice of this number was to ensure that the size of each portion of the input and prefix sum arrays stored by each processor was larger than the L3 cache, for all simulations. The size of the input array for  $10^9$  integers is 4 bytes/int  $\times 10^9$  integers = 4 GB. The size of the output arrays, which use 8 byte longs, is  $8 \times 10^9$  GB. In total, around 12 GB of memory is needed, which is well under the limit of 22 GB per node. The largest run of 256 cores still requires around 10 MB per core, which is greater than the size of the L3 cache.

### 3.2.2 Weak Scaling Study

A weak scaling study determines the efficiency of the algorithm as you increase the number of processors, while keeping the problem size *per core* fixed. The goal of a weak scaling study is to determine the increased cost of an algorithm as more processors are used to solve a bigger problem. The metric for the weak scaling studies used was efficiency, defined as  $E = T_p(n)/T_1(n) \times 100\%$ , noting that  $T_1$  is the time for the parallel algorithm with not processor, not the time for the serial algorithm. The ideal efficiency would be 100%, resulting in a flat line for  $E$  as a function of  $p$ .

For the weak scaling studies,  $10^8$  integers, per core, were used for both the MPI and OpenMP. This number was chosen for the same reasons as in the strong scaling study: the value is larger than the cache size and the total problem size will fit in the available RAM per node.

### 3.2.3 Determining Asymptotic Coefficients

To determine the validity of our theoretical analysis of the algorithms simulations were performed for a fixed number of processors with variable input sizes. This helps to determine if our theoretical model for run time as a function of input size and number of processors is sufficiently accurate. Our model does not account for communication latencies or memory access times. It also helps to determine in what range of problem sizes our asymptotic scalings are accurate.

The theoretical output time of the model is predicted as  $T_{pred}(t)$ . In general, the model  $T_{pred}(t)$  is a function of the problem size  $n$  and number of processors  $p$ . The model can be represented as  $T_{pred}(t) = C_0 g(n, p)$  for some  $n > n_0$ , for a given system. By performing experiments for various  $n$ , given a fixed  $p$ , the coefficients of the model  $C_0$  and  $n_0$  can be determined by plotting the ratio of the actual to predicted run times.

For both the OpenMP and MPI algorithm, the dominant term is expected to be  $O(n/p)$ . Rather than trying to fit a function to determine the leading coefficients in the model, which would be difficult due to statistical noise in the results and the small contribution from the  $O(p)$  terms, only the coefficients for the dominant  $O(n/p)$  term are determined. In the case of all algorithms, it is expected that the  $O(n/p)$  term will dominate.

## 3.3 Statistics

The experiments must be repeated to measure various forms of variability in the system, e.g., variable communication time, memory access times, inaccuracy of the timer, etc. Unless noted otherwise, the experiments were repeated 32 times for each specific simulations. The entire program is rerun for each iteration, to ensure the effect of variability in memory initialization times on execution time is represented accurately. simulations were performed to estimate the variability in the system. The entire simulation was repeated, to eliminate any unrealistic gains in efficiency due to.

From the 32 simulations, the reported run times are simply the average of the 32 simulations. The standard error in the average of a quantity is  $\sigma/\sqrt{N-1}$ , where  $\sigma$  is the standard deviation of the quantity. Since speed up and scalability are calculated quantities with a

statistical variance in both terms, it is necessary to approximate the error in the quantity. Based on the standard error propagation formula ([http://en.wikipedia.org/wiki/Propagation\\_of\\_uncertainty](http://en.wikipedia.org/wiki/Propagation_of_uncertainty)), the error for the ratio of two values is

$$\sigma\left(\frac{T_i}{T_j}\right) = \frac{T_i}{T_j} \sqrt{\left(\frac{\sigma_{T_i}}{T_i}\right)^2 + \left(\frac{\sigma_{T_j}}{T_j}\right)^2}. \quad (1)$$

The above equation is used to determine the standard error for all plotted speed ups and scalability. The plotted values are the 95% confidence interval, which is represented, assuming a Gaussian distribution of the error, as  $1.96\sigma$ .

## 4 Results

It is noted that although this algorithm in theory has the same (or slightly worse) time complexity as a linear array approach, there is less communication steps, which leads to overall a significantly more efficient algorithm as the cost of the communication steps becomes the limiting factor at higher MPI rank counts. This would have similar communication cost to a tree-traversal based algorithm.

### OPENMP

To limit memory access overhead times, it was necessary to allow threads to each have their own memory, rather than each accessing a portion of the array. So the algorithm is implemented slightly different than indicated. Otherwords the algorithm did not scale well.